An Open and Reproducible Edge-AI System for Litter Detection, Multiclass Classification, Tracking, and Counting

Systems and reproducibility preprint — verified software evidence and preliminary engineering validation

21 October 2026 | Abiola A. Babsalaam — Independent Researcher

⚠ AI & Data Notice: All output is AI-generated and must be independently verified before sharing with customers or taking action.

Abstract

Litter monitoring is often treated as an object-detection exercise, although a deployable solution must also ingest operational video, maintain stable class contracts, avoid repeated counting, persist evidence, support human review, and integrate safely with collection or sorting workflows. This paper presents an open and reproducible edge-AI system that addresses those requirements as one portable application. The software accepts attached cameras, images, recorded video, HTTP media, and RTSP streams. It normalizes public or custom detector output into a common schema, applies duplicate suppression and class-stabilized tracking, reports visible and unique-session counts, and stores run evidence in SQLite and JSON. A browser interface exposes annotated video, counts, review tasks, and source controls. The default operational taxonomy contains 16 litter categories and can map uncertain or hazardous items to human review and site-specific separation guidance. The detector is an interchangeable component rather than the product boundary: the current clean-clone workflow uses the public OmDet-Turbo Tiny open-vocabulary model, while a litter-specific RF-DETR checkpoint may be trained and published later if representative evaluation shows that adaptation is required. Preliminary engineering validation confirms installation, real inference, persistence, application tests, and model-only execution on NVIDIA CUDA and Apple silicon. On an AWS g4dn.xlarge with a Tesla T4, a repeated-image OmDet-Turbo benchmark measured 16.53 frames per second, 58.98 ms mean detector latency, 61.03 ms p95 latency, and 0.55 GiB peak allocated GPU memory. These measurements do not establish target-domain accuracy, full-pipeline sustained performance, or Jetson and AMD deployment performance. The paper therefore contributes a reusable system, explicit count semantics, a safety and security boundary, and a predeclared evaluation protocol while clearly separating verified software evidence from pending empirical claims.

Keywords: edge AI; litter detection; waste classification; object tracking; object counting; open-vocabulary detection; human-in-the-loop; NVIDIA Jetson; recycling; reproducible software

1. Introduction

Improperly discarded waste affects public-space quality, drainage, wildlife, and the movement of plastics into aquatic environments [1,2]. Collection teams typically depend on scheduled routes, visual inspection, public reports, and fixed infrastructure. These methods remain necessary, but they provide limited continuous information about where litter is present, what material it may contain, and whether an observed item was later collected.

Computer vision can provide an additional perception layer. A camera-equipped vehicle, fixed monitoring point, or sorting line can produce candidate object locations, predicted classes, confidence scores, and counts. Edge processing can keep raw imagery close to the source and continue operating when network connectivity is intermittent. However, a neural-network prediction alone is not an operational solution. The surrounding application must handle cameras and streams, normalize model output, preserve object identity across frames, define what a count means, retain auditable evidence, support corrections, and avoid unsafe coupling between uncertain perception and machinery.

Litter is also a difficult visual domain. Objects can be small, crushed, transparent, reflective, wet, dirty, partly occluded, or visually similar to permanent background objects. Material cannot always be inferred from appearance. A clear container, for example, may not contain enough visible evidence to distinguish glass from plastic. Public datasets demonstrate the feasibility of

litter detection but do not remove the need to validate the intended camera, environment, class ontology, and operating threshold [3–6].

This work treats the outcome—reliable detection, classification, tracking, counting, review, and integration—as the unit of design. It does not compare model families or attempt to establish a leaderboard. The detector is placed behind a stable adapter so the complete workflow can operate with a publicly available model and later accept a validated site-specific checkpoint without changing downstream data contracts.

The result is the open-source Edge Litter Detection application [7]. A user can clone the repository and run one launcher that prepares an isolated environment, installs dependencies, downloads a pinned public model, discovers or accepts a source, performs inference, and writes durable evidence. The same application supports a local dashboard, a private HTTP API, SQLite queries, human review, and integration events for a larger vehicle or facility system.

This paper is released as a systems and reproducibility preprint. It documents the implemented design and preliminary technical evidence while the target-domain dataset and target-device acceptance tests are pending. This distinction is deliberate: model detections on demonstration media are functional evidence, not accuracy evidence, and a model-only accelerator benchmark is not proof of sustained camera performance.

1.1 Research questions

The work is organized around the following questions:

1. How can litter detection, multiclass classification, tracking, counting,

storage, review, and integration be packaged as a reusable edge application?

2. Can the application operate immediately with a public model, without

requiring every deployment user to collect data or train a network?

3. How should visible observations, unique tracked objects, and confirmed

collections be represented so that operational reports do not overstate the outcome?

4. What reproducible protocol is required to evaluate detection accuracy,

counting accuracy, throughput, memory, power, thermal behavior, and stability on the intended device?

5. What privacy, security, and machinery-safety boundaries are necessary when

the system is used on a cleanup vehicle or sorting line?

1.2 Contributions

The principal contributions are:

- a portable, model-adapter-based application for image, camera, video, HTTP, and RTSP litter detection;
- a 16-class operational taxonomy with configurable prompts and separation destinations;
- tracking-based visible and unique-session counts, with independently confirmed collections stored as a separate quantity;

- an offline-first SQLite evidence model, review queue, browser dashboard, API, and durable integration outbox;
- one-command installation and execution using public, pinned model artifacts;
- deployment paths for CPU, NVIDIA CUDA, NVIDIA Jetson, AMD ROCm, and Apple silicon;
- model, data, privacy, security, and safety boundaries intended for reuse in real collection and sorting workflows; and
- a predeclared empirical protocol that prevents preliminary engineering observations from being reported as final accuracy or target-device results.

2. Related Work

2.1 Litter and waste datasets

TACO introduced a publicly available dataset of waste “in the wild,” with object annotations and a detailed material-oriented taxonomy [3]. It is useful for bootstrapping because it contains varied contexts rather than only controlled conveyor imagery. Studies based on TACO have also highlighted the difficulty of small objects, class imbalance, and the risk of reporting one aggregate metric without examining precision, recall, F1, object size, and scene conditions [4,5].

ZeroWaste focuses on industrial waste sorting and emphasizes dense, overlapping objects and material recovery facility conditions [6]. It illustrates why a sorting-line camera and a road-facing vehicle camera should not automatically be treated as the same domain. More recent datasets such as RoLID-11K expand real-world litter coverage and underline continuing interest in diverse scenes and fine-grained categories [8]. Dataset growth is valuable, but deployment performance still depends on the intended viewpoint, camera, lighting, object size, contamination, motion, and taxonomy.

The present project therefore uses public data as attributed bootstrap material rather than as proof of field performance. The repository can prepare a grouped TACO-derived COCO dataset while retaining source and license metadata. Target-domain samples remain necessary before a litter-specific checkpoint or an operational accuracy claim can be released.

2.2 Detection and open-vocabulary models

Modern object detectors predict class labels and bounding boxes, allowing multiple litter items to be localized within one image. COCO established widely used detection evaluation conventions, including average precision across intersection-over-union thresholds [9]. DETR reformulated detection as direct set prediction with transformer-based global reasoning [10], and DINO subsequently improved DETR-style training and query formulation [11].

Grounding DINO extends detection toward open-set, text-conditioned operation [12]. OmDet-Turbo is designed as a real-time open-vocabulary detector that can process user-defined class names [13]. These models are useful where a public checkpoint must work before a task-specific training set exists. They also introduce prompt sensitivity and semantic ambiguity: a text-conditioned model can return a plausible box and class without being calibrated for the local waste stream.

RF-DETR is a recent real-time detection-transformer family with deployment-oriented model sizes and training support [14]. In this project it is an optional path for a centrally trained, fixed-taxonomy checkpoint if the public open-vocabulary path does not satisfy the use-case gates. This is not a comparison between OmDet-Turbo and RF-DETR. OmDet-Turbo currently enables immediate use and data collection; a custom checkpoint is a future adaptation mechanism, not a requirement imposed on ordinary users.

2.3 Tracking and counting

Video detection creates repeated observations of the same object. Summing all frame detections can severely overcount persistent litter. Tracking methods associate detections over time and assign identities. SORT demonstrated a simple online combination of motion estimation and data association [15]. ByteTrack showed that lower-confidence detections can improve association when handled appropriately [16].

The current application uses a compact class-stabilized intersection-over-union tracker to minimize runtime dependencies. It is suited to functional use with fixed or slowly moving cameras and mostly stationary litter. It is not presented as a universal identity solution. Fast camera motion, long occlusion, scene cuts, loops, and crowded items can create duplicate or merged tracks. Accordingly, count error must be evaluated separately from detection average precision, using track-annotated video.

2.4 Edge deployment and reproducibility

Edge deployment adds preprocessing, decoding, memory transfer, visualization, storage, and synchronization to model inference. Performance depends on the complete software and hardware condition, including input size, precision, runtime, power mode, cooling, camera pipeline, and co-resident workload. PyTorch, Transformers, and OpenCV provide portable implementation components [17–19], but portability does not imply equal performance on all devices.

The reproducibility gap is often caused by missing model revisions, thresholds, environment versions, hardware details, or test-media checksums. This project records run metadata and emits machine-readable evidence so results can be connected to a source, model, configuration, device, and Git revision.

3. Requirements and Use Cases

The software is intended for four related uses:

1. A vehicle-mounted camera observes roadside litter and associates candidate

objects with a device, vehicle, and route.

2. A fixed camera monitors a location or accumulation point.

3. A recycling or waste-sorting line detects visible items and presents a

suggested class and operator-configured destination.

4. A human worker uses the annotated feed and review queue to find, classify,

or separate uncertain items.

The application is a perception and evidence component. It is not an autonomous enforcement tool and is not a safety-rated motion controller. Person identification, face recognition, license-plate recognition, behavior scoring, and automated penalties are outside the project scope.

3.1 Functional requirements

The system must:

- accept a local camera index, image, video, HTTP resource, or RTSP stream;
- produce zero or more boxes, class names, confidence scores, and timestamps;
- preserve a stable class ID and name contract downstream of the detector;
- reduce duplicate boxes and associate video observations into tracks;
- report current visible counts and run-scoped unique counts by class;
- persist run, frame, count, review, and integration evidence locally;
- remain useful without continuous cloud connectivity;
- allow a human to confirm, correct, or reject uncertain classifications;
- store actual collection confirmations separately from observations;
- expose bounded local controls through a dashboard and API; and
- retain enough environment and revision data to reproduce a reported run.

3.2 Non-functional requirements

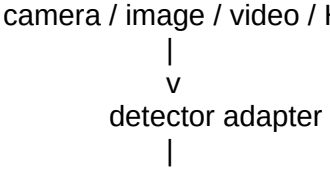
The system must:

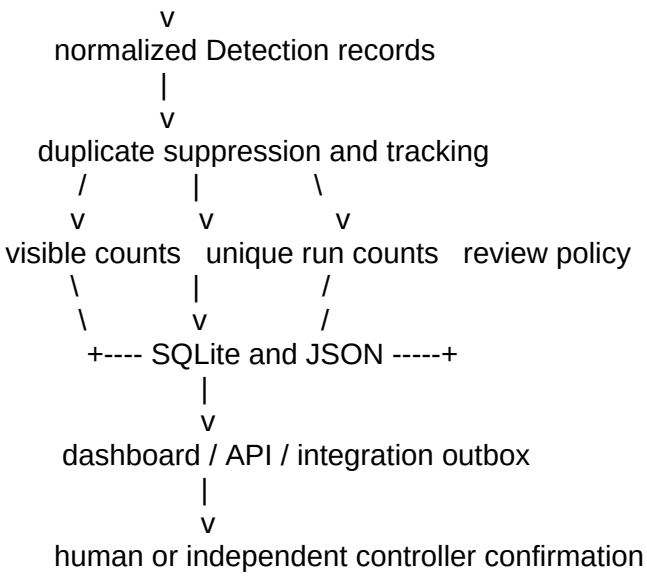
- install from a clean clone without private credentials or a private model;
- use an isolated Python environment and avoid machine-specific paths;
- function on 64-bit x86_64 and ARM64 systems;
- avoid dependence on one accelerator vendor;
- default to local-only network exposure;
- validate and bound uploaded media and remote-source behavior;
- support model artifact pinning and integrity checks;
- avoid direct machinery actuation;
- limit unnecessary retention of raw media and identifiers; and
- disclose when evidence is preliminary, incomplete, or device-specific.

4. System Design

4.1 End-to-end architecture

The runtime flow is:





Capture, detection, counting, persistence, and integration are separate components. A detector adapter converts backend-specific results into a stable record containing class ID, class name, confidence, and bounding box. This allows the rest of the application to remain unchanged when a model or accelerator path is replaced.

The implementation currently includes Grounding DINO Tiny, OmDet-Turbo Tiny, and RF-DETR adapters. OmDet-Turbo Tiny is the recommended no-training path in the documented demonstration because it uses a public checkpoint and returns confidence-scored open-vocabulary detections. Grounding DINO remains a functional bootstrap and review aid. RF-DETR accepts a future project-maintained or site-specific checkpoint. No backend is described as accurate for a target site until it has passed the same frozen use-case evaluation.

4.2 Operational taxonomy

The default taxonomy contains 16 categories:

ID	Class	ID	Class
0	plastic_bottle	8	cigarette_butt
1	glass_bottle	9	food_waste
2	metal_can	10	textile
3	paper	11	tire
4	cardboard	12	battery
5	plastic_bag	13	electronic_waste
6	food_wrapper	14	broken_glass
7	disposable_cup	15	other_litter

This taxonomy balances common item types, material-separation needs, hazardous items, and an explicit ambiguous class. Public-model prompts are configurable without changing the downstream class contract. Separation destinations are kept in a different configuration because recycling policy and facility streams vary by operator and jurisdiction.

The taxonomy is not a claim that every material can be visually identified. Low-confidence, hazardous, and ambiguous observations can enter a human review queue. If a released closed-set checkpoint is trained later, every advertised class must have documented training and held-out evaluation support.

4.3 Tracking and count semantics

The system deliberately reports different quantities:

- visible_counts are detections present in the current frame;
- confirmed_visible_counts are currently visible detections whose tracks

have met a confirmation threshold;

- unique_session_counts are confirmed track identities counted once in one

run; and

- confirmed_collected_counts are physical pickups separately confirmed by a

human, collection controller, sensor-fusion service, or reviewed import.

These quantities answer different questions. Visible counts estimate the current scene. Unique-session counts estimate how many track identities were encountered during a run. Confirmed-collected counts represent an external outcome. A detection does not prove collection, and unique counts from different runs cannot be added without accepting possible duplicate physical objects.

The default tracker associates boxes by overlap, retains a track across short missed intervals, confirms it after repeated observations, and stabilizes the class using accumulated votes. This reduces frame-by-frame overcounting and label flicker while preserving a dependency-light edge path. For fast-moving vehicle footage or crowded sorting lines, the same detection and event schema can support a stronger tracker or geospatial deduplication after separate validation.

4.4 Persistence and offline operation

The application writes streaming JSON Lines events, an end-of-run JSON summary, optional annotated media, and a SQLite operational database. SQLite uses write-ahead logging and parameterized statements. The schema includes:

Table	Purpose
runs	source, device, vehicle, route, status, timing, and final counts
frame_observations	per-frame visible and cumulative unique counts
detections	optional detailed boxes and tracks
run_class_counts	query-friendly final class totals
review_queue	pending, confirmed, corrected, or rejected classifications
collection_confirmations	independently confirmed physical pickups
integration_outbox	durable events waiting for an external acknowledgement

Detailed detection rows are disabled by default to reduce long-term storage. Operators can retain counts and summaries while applying a shorter retention period to raw or annotated media. The outbox pattern permits a vehicle to continue operating without network connectivity and forward events later through an organization-controlled authenticated transport.

4.5 Human-assisted separation

Predicted classes can be mapped to site-specific material streams. A dashboard or external application can display the class, confidence, destination, and review state to a worker. Corrections are auditable and do not silently rewrite the original model observation.

The separation mapping is advisory. Composite, contaminated, hazardous, or unknown material may require a manual decision. Batteries, electronic waste, broken glass, sharps, chemicals, and pressurized containers require local safety procedures beyond a visual prediction.

4.6 Dashboard and API

The FastAPI service provides health and readiness endpoints, image detection, stream-session control, count queries, review operations, collection confirmation, and outbox acknowledgement. It also serves a dependency-free browser dashboard. A server-side session can open a camera or file on the edge computer, avoiding the false assumption that a browser camera index refers to the remote device.

The API binds to loopback by default and has no built-in multi-user identity system. Remote access should use an SSH tunnel or a mature authenticated gateway. The project does not recommend exposing the service directly to the public Internet.

4.7 Machinery-safety boundary

Perception output may inform an operator or task planner, but it must not directly move hydraulic, pneumatic, or electric equipment. A production collection arm requires an independent safety-rated controller, emergency stops, manual override, exclusion-zone and proximity sensing, arm-position and load limits, stale-signal handling, and positive pickup confirmation. If safety inputs are missing or contradictory, motion must remain inhibited.

This boundary prevents a probabilistic class or box from becoming an unreviewed actuation command. It also preserves the difference between “an object was observed” and “an object was collected.”

5. Implementation and Reproducibility

5.1 Software release

The implementation described in this preprint is Edge Litter Detection version 0.10.1 at Git commit 0624677b72cf1d16b59a335b78d9003f2b479786 [7]. Code and project documentation are released under Apache License 2.0. Model checkpoints and datasets retain their own licenses and are not relicensed by the repository.

The primary entry point is:

```
git clone https://github.com/biolababs/edge-litter-detection.git
cd edge-litter-detection
./startup.sh --backend omdet_turbo --demo-video --no-display
```

The launcher creates or reuses a virtual environment, installs the selected runtime, downloads attributed demonstration media when requested, downloads a pinned public checkpoint, runs inference, and writes evidence. Attached cameras and network streams use the same path:

```
./startup.sh --source 0 --backend omdet_turbo

export LITTER_STREAM_URL="rtsp://camera-host/live"
./startup.sh --source "$LITTER_STREAM_URL" --backend omdet_turbo --no-display
```

The dashboard is started with:

```
./startup.sh --dashboard --demo-video --backend omdet_turbo --device auto
```

Counts remain queryable after the process exits:

```
.venv/bin/python query_counts.py runs --limit 1
.venv/bin/python query_counts.py summary
```

5.2 Model artifact controls

Public Hugging Face model revisions are immutable in the application configuration, SafeTensors is preferred, and remote model code is disabled. Custom remote model references must be pinned to a commit and accompanied by a checksum where the download path requires it. Demonstration downloads use HTTPS, redirect validation, size limits, media-signature checks, temporary files, and SHA-256 records.

The repository keeps model selection separate from model availability. A public model makes the application immediately deployable for functional use, but only a checkpoint that passes representative accuracy and systems gates should be described as accepted for an operational site.

5.3 Hardware portability

The minimum functional baseline is a 64-bit x86_64 or ARM64 computer with four logical CPU cores, 8 GiB of memory, 10 GiB of free storage, Python 3.10–3.14, and network access for initial installation. An accelerator is optional for installation checks, still images, API integration, and low-rate processing.

The recommended real-time evaluation baseline is six or more modern logical cores, 16 GiB host memory for a discrete-GPU system, 8 GiB or more accelerator or unified memory, 20 GiB free SSD storage, and a vendor-supported CUDA, ROCm, or JetPack PyTorch environment. These values are capacity-planning baselines, not guaranteed performance.

The code exposes CPU, NVIDIA CUDA, AMD ROCm through PyTorch, Apple MPS, and Jetson-oriented paths. Each target must retain its own environment and benchmark record. Results measured on a Tesla T4 cannot be presented as Jetson, AMD, or Apple results.

5.4 Software verification and secure development

At the reported revision, 73 automated tests cover configuration, model adapters, tracking, storage, streaming, benchmark output, training-data preparation, evaluation, and API security

behavior. The release workflow also checks command-line help paths, shell syntax, package structure, startup regression, static correctness, dependency vulnerabilities, and Python package construction.

The security design is organized around concepts from the NIST Secure Software Development Framework and OWASP Application Security Verification Standard [20,21]. Controls include loopback network defaults, bounded media uploads, content-signature validation, browser security headers, parameterized SQL, non-root containers, pinned automation, automated dependency updates, secret scanning, CodeQL analysis, and a private vulnerability-reporting path. This is a secure-development review, not a penetration test, certification, or claim of formal compliance.

6. Evaluation Protocol

6.1 Evaluation principle

The evaluated object is one frozen solution configuration: checkpoint, taxonomy, prompts or class head, thresholds, image size, tracker settings, runtime, and Git revision. The purpose is to determine whether the complete system satisfies the use case. The study does not compare detectors for rank.

6.2 Data acquisition and governance

Evaluation data should contain the intended truck, sorting-line, fixed-camera, or handheld viewpoint; expected object sizes; empty scenes; confusing background objects; lighting and weather variation; motion blur; occlusion; crushing; contamination; and class imbalance. Sources, licenses, collection dates, sites, sessions, and redistribution restrictions must be recorded.

Adjacent frames and repeated views from one video are strongly correlated. Random image-level splitting can place nearly identical content in development and test partitions. Entire source videos, routes, sites, or collection sessions should therefore remain in only one partition. The held-out test set must be frozen and checksummed before checkpoint, threshold, or prompt selection is finalized.

Where feasible, a quality-control subset should be labeled independently by at least two annotators. The annotation guide must define partial occlusion, truncation, minimum visible evidence, object grouping, ambiguous material, and non-litter exclusions. Identifiable people, plates, private property, and unnecessary location metadata should be removed or de-identified under the approved protocol.

6.3 Detection and classification measures

The final evaluation will report:

- precision, recall, and F1 at the frozen operational threshold;
- mAP50 and mAP50–95;
- per-class average precision and support;
- a confusion analysis;
- results by object size, lighting, weather, clutter, occlusion, distance, and

scene type; and

- human-review rate and correction rate.

No class should be advertised as supported if it lacks held-out examples. Aggregate mAP must not conceal poor recall for a hazardous or operationally important minority class.

6.4 Counting measures

Track-annotated video is required for identity-based counting claims. The final evaluation will report per-frame absolute error, visible-count mean absolute error, unique-session count error, duplicate-track rate, missed-track rate, identity switches where ground truth permits, and class reassignment within a track. Results should be stratified by camera motion, object density, occlusion, scene cuts, and looped or repeated routes.

Detection accuracy is necessary but insufficient for count accuracy. A detector can have acceptable image metrics while an unstable tracker repeatedly counts the same item.

6.5 Systems measures

Systems evaluation will separate:

- pure detector-call latency;
- preprocessing, inference, and postprocessing latency;
- observed source-to-result throughput;
- first-inference and steady-state latency;
- p50 and p95 latency;
- process, system, and accelerator memory;
- power and energy per processed frame where supported;
- temperature, clocks, and throttling state; and
- sustained-run failures, memory growth, camera interruptions, and restart behavior.

Warm-up must precede repeated samples. Principal target-device conditions should be repeated at least three times. Input resolution, precision, power mode, cooling, camera or media source, co-resident workload, software versions, and optimized artifact checksum must remain recorded with the result.

6.6 Target-device acceptance

The provisional engineering target is at least 15 source-to-result frames per second for the declared camera and coexistence condition. For the shared Jetson Orin Nano 8 GB, the preferred incremental detector memory is no more than 1.5 GiB, with a 2.0 GiB initial hard limit and at least 2.0 GiB remaining available memory. These are acceptance criteria, not achieved results.

An optimized FP16 or INT8 artifact must be checked against the same held-out data. INT8 requires representative calibration samples and should not be accepted merely because it is faster.

7. Preliminary Validation

7.1 Scope

Preliminary validation establishes that the software can be installed, executed, tested, and measured. It does not establish target-domain detection accuracy, counting accuracy, operational benefit, environmental impact, or performance on every supported hardware family.

7.2 Functional software evidence

The clean-clone workflow has completed real OmDet-Turbo image inference through the normal launcher, produced normalized detections, written JSON and annotated output, and persisted run and count data in SQLite. The dashboard, API, tracking, review, storage, benchmark, and preparation paths are covered by 73 automated tests at the reported revision.

One TACO demonstration image produced six unique-session detections in a single-image run: three metal_can, one plastic_bag, one plastic_bottle, and one tire. These labels are reported only to demonstrate the pipeline and record schema. They are not an accuracy result because the image was not evaluated under the frozen taxonomy mapping and operating threshold, and some predictions do not match the source annotations.

7.3 Preliminary model-only measurements

Platform	Backend	Samples	Throughput	Mean detector latency	p95 latency	Accelerator memory
AWS g4dn.xlarge, Tesla T4	OmDet-Turbo Tiny, CUDA	50 repeated images after 10 warm-up calls	16.53 FPS	58.98 ms	61.03 ms	0.55 GiB peak allocated
Apple silicon, MPS	OmDet-Turbo Tiny	10 repeated images after 3 warm-up calls	5.97 FPS	160.02 ms	164.24 ms	not available

The T4 measurement used a repeated image and measured the detector benchmark path rather than sustained video capture, annotation rendering, database writing, and dashboard delivery. The original output was produced on the AWS instance; the public repository currently retains a transcribed evidence record and the documented values, not the original remote JSON file. The Apple measurement is retained as a repository JSON artifact. Neither result measures accuracy.

The T4 result demonstrates that the current public-model path can exceed the provisional 15 FPS engineering target in one discrete-GPU, model-only condition. It must not be generalized to an 8 GB Jetson, an AMD device, or the full application.

7.4 Read-only Jetson inventory

A read-only inventory identified the intended target as a Jetson Orin Nano Engineering Reference Developer Kit Super with Ubuntu 22.04.5, JetPack 6.2.3, L4T 36.5.2, TensorRT

10.3.0, Python 3.10.12, and 7.4 GiB physical memory. At capture time, 4.5 GiB was available, no /dev/video* camera was present, and active local LLM, application, and development services were observed.

No package was installed, no service was stopped, and no configuration was changed. The inventory confirms platform constraints only. Jetson inference, accuracy, throughput, memory, power, thermal, and sustained coexistence tests remain pending an authorized test window.

7.5 Evidence status

Claim area	Status
Clean clone, launcher, public-model inference, SQLite, dashboard/API	implemented and functionally verified
Automated software tests and release/security workflows	verified at reported revision
Tesla T4 OmDet-Turbo model-only execution	preliminary measured evidence
Apple MPS OmDet-Turbo model-only execution	preliminary measured evidence
Target-domain detection/classification accuracy	pending representative labeled data
Video counting accuracy	pending track-annotated video
Jetson Orin 8 GB inference, power, thermals, and sustained stability	pending authorized test
AMD ROCm / Strix Halo inference and sustained performance	pending target-device test
Cleanup or sorting operational benefit	pending field evaluation

8. Discussion

The implementation shows why an outcome-focused litter system should not be defined by a detector name. Public open-vocabulary detection makes the application usable before a local training project exists. The same runtime then creates review data and operational evidence that can justify—or avoid—the cost of custom training. This lowers the entry barrier for organizations that need to test the workflow but do not have an annotation team.

The adapter boundary also prevents downstream lock-in. A truck or facility integration consumes stable detections, counts, reviews, and collection confirmations rather than a backend-specific tensor. If a future checkpoint is accepted, the storage and integration contracts remain unchanged.

Explicit count semantics are a central design result. “Detected,” “uniquely tracked,” and “collected” are not synonyms. A useful operational dashboard must show which quantity it is presenting and must preserve uncertainty. Otherwise, a system can report impressive totals while repeatedly observing the same item or assuming that a proposed pickup succeeded.

The preliminary T4 measurement supports continued evaluation of OmDet-Turbo as the no-training path. Its 16.53 FPS detector-only throughput in the measured condition is encouraging, but the gap between that number and a full, sustained camera workload can be material. The planned protocol therefore records capture, visualization, storage, thermal, and co-resident workload effects instead of treating model latency as an end-to-end result.

The Jetson inventory demonstrates another practical constraint: an edge device may already host valuable services. A benchmark that stops the existing workload or consumes swap can answer a laboratory question while failing the deployment requirement. Coexistence must be part of the declared condition, and exclusive benchmarking should be labeled separately if it is later authorized.

For human-assisted sorting, confidence and review state may be as important as the top class. The system can route ambiguous predictions to an operator instead of forcing automation. This supports a graduated deployment: visual assistance first, operational measurement second, and only then any separately safety-engineered automation.

9. Limitations and Threats to Validity

First, the public open-vocabulary checkpoint has not been calibrated on a representative target-domain dataset. Demonstration detections may be useful for software testing while being inaccurate for operational decisions.

Second, the preliminary benchmarks use repeated images. They reduce source variation and do not include all full-pipeline costs. Sample counts are modest, and the remote T4 JSON artifact was not copied back into the repository before the instance output was summarized.

Third, the built-in tracker is intentionally compact. It may duplicate, merge, or lose identities under fast camera motion, long occlusion, dense overlap, abrupt scene cuts, or repeated routes. Counting claims require track-level ground truth and may require a stronger tracker or spatial deduplication.

Fourth, the taxonomy contains visually ambiguous material categories. Contamination, composite packaging, transparent objects, and damage may make a single correct visual class impossible. Human review and an other_litter outcome mitigate but do not eliminate this limitation.

Fifth, Jetson and AMD execution paths are documented but not yet empirically validated. The paper does not claim that the measured T4 rate will transfer to an 8 GB unified-memory platform.

Sixth, no field study has yet measured whether the system improves route planning, worker productivity, recovery rate, separation quality, or cleanup outcomes. Those benefits remain hypotheses for later operational evaluation.

Finally, the security review is code- and configuration-oriented. A production deployment still requires environment-specific authentication, authorization, network segmentation, logging, patching, retention, privacy assessment, and safety engineering.

10. Ethics, Privacy, Security, and Sustainability

The application detects litter rather than identities. Cameras should be positioned to minimize people, private property, and license plates. Raw video retention should be limited, access controlled, and separated from public research artifacts. Public-space recording must follow the relevant consent, notice, ethics, labor, and data-protection requirements before collection.

Model performance may vary across communities, environments, seasons, and camera quality. Litter observations should not be used as an unreviewed proxy for individual behavior,

neighborhood worth, or enforcement priority. Error analysis should examine locations and conditions where recall or false positives differ.

The local API is intentionally private by default. Production operators must add an authenticated access layer, restrict outbound stream access, protect credentials, and maintain supported dependencies. Model and media downloads should remain pinned and verified.

Edge processing can reduce continuous raw-video transmission but still consumes electricity and hardware. Final evaluation should report power and energy per processed frame where feasible and should select the least complex configuration that satisfies the use case. Environmental benefit should not be inferred from detection metrics alone.

11. Data, Code, and Reproducibility Availability

The software, configuration, documentation, tests, and reproduction commands are available at:

https://github.com/biolababs/edge-litter-detection

This paper describes version 0.10.1 at commit:

0624677b72cf1d16b59a335b78d9003f2b479786

The repository includes attributed TACO demonstration preparation and a license-aware sorting-floor demonstration downloader. Downloaded media, third-party datasets, and model checkpoints retain their original terms and are not redistributed under the code license.

Preliminary evidence records accompany this manuscript under paper/evidence/. Runtime-generated records normally appear under reports/generated/, and the SQLite store appears under state/litter-edge.db. Transient media, private data, credentials, and raw device inventories are excluded from version control.

12. Conclusion and Future Work

This paper presents a reusable edge-AI system for litter detection, multiclass classification, tracking, counting, review, and integration. The main contribution is the complete operational boundary rather than a new detector. A public open-vocabulary model permits immediate clean-clone use, while the adapter design preserves a path to a centrally trained fixed-taxonomy checkpoint when representative evidence requires it.

The implementation distinguishes visible observations, unique tracked objects, and independently confirmed collections; stores durable offline evidence; supports human correction; and prevents direct perception-to-actuation coupling. Preliminary validation confirms a functioning codebase and provides bounded T4 and Apple-silicon model-only measurements. It does not yet support claims of target-domain accuracy or Jetson/AMD production performance.

The next empirical work is to collect and annotate representative field and sorting footage, freeze grouped development and test partitions, establish inter-annotator quality, select and freeze one operating configuration, evaluate per-class detection and count error, and complete authorized sustained tests on Jetson Orin 8 GB and AMD hardware. A subsequent paper revision should replace pending evidence with traceable accuracy, power, thermal, reliability, and operational-outcome results.

Declarations

Author contributions

Abiola A. Babsalaam conceived the use case, designed and implemented the software, prepared the preliminary validation, and authored the manuscript.

Funding

No external funding is declared for this preprint.

Competing interests

The author declares no competing interests.

Ethics statement

This preprint reports software development, public demonstration media, and technical benchmarks. It does not report a study involving recruited human participants. Any future field collection containing identifiable people, vehicles, private property, or precise locations must receive the required institutional, organizational, and legal approval before collection or publication.

Copyright and licensing

Repository code and project documentation are available under Apache License 2.0. Third-party datasets and model artifacts retain their original licenses. Copyright in this manuscript remains with the author unless a separate publication license is selected.

References

[1] Kaza, S., Yao, L. C., Bhada-Tata, P., & Van Woerden, F. (2018). What a Waste 2.0: A Global Snapshot of Solid Waste Management to 2050. World Bank. https://doi.org/10.1596/978-1-4648-1329-0

[2] Jambeck, J. R., Geyer, R., Wilcox, C., Siegler, T. R., Perryman, M., Andrady, A., Narayan, R., & Law, K. L. (2015). Plastic waste inputs from land into the ocean. Science, 347(6223), 768–771. https://doi.org/10.1126/science.1260352

[3] Proença, P. F., & Simões, P. (2020). TACO: Trash Annotations in Context for Litter Detection. arXiv:2003.06975. https://doi.org/10.48550/arXiv.2003.06975

[4] Córdova, M., Pinto, A., Hellevik, C. C., Alaliyat, S. A.-A., Hameed, I. A., Pedrini, H., & Torres, R. d. S. (2022). Litter detection with deep learning: A comparative study. Sensors, 22(2), 548. https://doi.org/10.3390/s22020548

[5] Rad, M. S., von Kaenel, A., Droux, A., Tièche, F., Ouerhani, N., Ekenel, H. K., & Thiran, J.-P. (2017). A computer vision system to localize and classify wastes on the streets. In Computer Vision Systems (pp. 195–204). https://doi.org/10.1007/978-3-319-68345-4_18

[6] Bashkirova, D., Abdelfattah, M., Zhu, Z., Akl, J., Alladkani, F., Hu, P., Ablavsky, V., Calli, B., Adel Bargal, S., & Saenko, K. (2022). ZeroWaste dataset: Towards deformable object segmentation in cluttered scenes. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition. https://doi.org/10.48550/arXiv.2106.02740

[7] Babsalaam, A. A. (2026). Edge Litter Detection, Multiclass Classification, Tracking, and Counting (Version 0.10.1) [Computer software]. GitHub. https://github.com/biolababs/edge-litter-detection

[8] Wu, T., Xu, Q., He, X., Weekes, O., Brown, J., & Duan, W. (2026). RoLID-11K: A dashcam dataset for small-object roadside litter detection. arXiv:2601.00398. https://doi.org/10.48550/arXiv.2601.00398

[9] Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., & Zitnick, C. L. (2014). Microsoft COCO: Common objects in context. In Computer Vision – ECCV 2014 (pp. 740–755). https://doi.org/10.1007/978-3-319-10602-1_48

[10] Carion, N., Massa, F., Synnaeve, G., Usunier, N., Kirillov, A., & Zagoruyko, S. (2020). End-to-end object detection with transformers. In Computer Vision – ECCV 2020 (pp. 213–229). https://doi.org/10.1007/978-3-030-58452-8_13

[11] Zhang, H., Li, F., Liu, S., Zhang, L., Su, H., Zhu, J., Ni, L. M., & Shum, H.-Y. (2022). DINO: DETR with improved denoising anchor boxes for end-to-end object detection. arXiv:2203.03605. https://doi.org/10.48550/arXiv.2203.03605

[12] Liu, S., Zeng, Z., Ren, T., Li, F., Zhang, H., Yang, J., Jiang, Q., Li, C., Yang, J., Su, H., Zhu, J., & Zhang, L. (2023). Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection. arXiv:2303.05499. https://doi.org/10.48550/arXiv.2303.05499

[13] Zhao, T., Liu, P., He, X., Zhang, L., & Lee, K. (2024). Real-time transformer-based open-vocabulary detection with efficient fusion head. arXiv:2403.06892. https://doi.org/10.48550/arXiv.2403.06892

[14] Robinson, I., Robicheaux, P., Popov, M., Ramanan, D., & Peri, N. (2025). RF-DETR: Neural architecture search for real-time detection transformers. arXiv:2511.09554. https://doi.org/10.48550/arXiv.2511.09554

[15] Bewley, A., Ge, Z., Ott, L., Ramos, F., & Upcroft, B. (2016). Simple online and realtime tracking. In 2016 IEEE International Conference on Image Processing (pp. 3464–3468). https://doi.org/10.1109/ICIP.2016.7533003

[16] Zhang, Y., Sun, P., Jiang, Y., Yu, D., Weng, F., Yuan, Z., Luo, P., Liu, W., & Wang, X. (2022). ByteTrack: Multi-object tracking by associating every detection box. In Computer Vision – ECCV 2022 (pp. 1–21). https://doi.org/10.1007/978-3-031-20047-2_1

[17] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. In Advances in Neural Information Processing Systems 32.

[18] Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M., et al. (2020). Transformers: State-of-the-art natural language processing. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations (pp. 38–45). https://doi.org/10.18653/v1/2020.emnlp-demos.6

[19] Bradski, G. (2000). The OpenCV library. Dr. Dobb's Journal of Software Tools.

[20] National Institute of Standards and Technology. (2022). Secure Software Development Framework (SSDF) Version 1.1: Recommendations for Mitigating the Risk of Software Vulnerabilities (NIST SP 800-218). https://doi.org/10.6028/NIST.SP.800-218

[21] OWASP Foundation. (2021). Application Security Verification Standard 4.0.3. https://owasp.org/www-project-application-security-verification-standard/

Appendix A. Reproduction Commands

A.1 Functional demonstration

```
git clone https://github.com/biolababs/edge-litter-detection.git
cd edge-litter-detection
git checkout 0624677b72cf1d16b59a335b78d9003f2b479786
./startup.sh --backend omdet_turbo --demo-video --no-display
```

A.2 Dashboard

```
./startup.sh --dashboard --demo-video --backend omdet_turbo --device auto
```

Open http://127.0.0.1:8088. For a remote edge computer, keep the service on loopback and use an SSH tunnel.

A.3 Counts

```
.venv/bin/python query_counts.py runs --limit 10
.venv/bin/python query_counts.py summary
.venv/bin/python query_counts.py reviews --status pending
```

A.4 Benchmark

```
.venv/bin/python benchmark.py \
  --backend omdet_turbo \
  --source media/taco-demo/taco-000866.jpg \
  --device cuda \
  --warmup 10 \
  --image-repeats 50 \
  --max-samples 50 \
  --output reports/generated/omdet-benchmark.json
```

A.5 Validation

```
python -m pip install -e ".[dev]"
make validate
make security
```

Appendix B. Required Evidence for the Empirical Revision

- approved data-governance, privacy, and annotation protocol;
- source-grouped and checksummed development and held-out test partitions;
- support counts for every released class;
- checkpoint or public-model revision, prompt/class schema, thresholds, and configuration checksum;
- detection, classification, and counting accuracy with per-class results;
- target-device environment manifest;
- full-pipeline latency, throughput, memory, power, temperature, and sustained stability;
- field or sorting-line operational outcomes; and
- updated limitations tied to measured failure modes.

Editorial Transparency and Verification Notice

⚠ AI & Data Notice: All output is AI-generated and must be independently verified before sharing with customers or taking action.